

Comparative Study of Generative Models on CIFAR-10

Archil Zhghenit, Nika Matcharadze

Free Uni of Tbilisi

January 31, 2026

CIFAR-10

- 60,000 color images (32×32)
- 10 classes (airplane, automobile, bird, cat, etc.)
- 50k training / 10k test images
- Images normalized to $[-1, 1]$ (or $[0, 1]$)

ConvVAE

- Convolutional Resnet Style encoder–decoder architecture
- Gaussian latent space with KL divergence regularization
- Optimizes ELBO objective

ConvVAE (β -VAE Variant)

- Convolutional encoder–decoder with Gaussian latent variables
- Objective:

$$\mathcal{L} = \mathbb{E}_{q(z|x)}[\log p(x|z)] - \beta \text{KL}(q(z|x) \| p(z))$$

- β controls the trade-off between reconstruction fidelity and latent regularization

Effect of β in ConvVAE

β	FID ↓	IS ↑	Mutual Information
0.1	129	3.66	5.54
1	160	2.00	5.53
10	313	1.70	0.00


Analysis: ConvVAE with Different β

- $\beta = 0.1$: Best generative performance (lowest FID, highest IS)
- $\beta = 1$: Stronger regularization with moderate quality degradation
- $\beta = 10$: Posterior collapse observed ($MI \approx 0$)
- High β harms sample quality on CIFAR-10 despite stronger constraints

Generated Samples (ConvVAE)



$\beta = 0.1$




$\beta = 1$



$\beta = 10$


Reconstruction Samples (ConvVAE)



$\beta = 0.1$



$\beta = 1$



$\beta = 10$

MMD-VAE

- Replaces KL divergence with Maximum Mean Discrepancy (MMD)
- Encourages richer latent representations
- Better sample diversity than standard VAE

Evolution of the VAE Loss

- **Vanilla VAE Loss (ELBO)**

$$\mathcal{L}_{\text{VAE}} = \mathbb{E}_{q(z|x)}[\log p(x|z)] - \text{KL}(q(z|x) \parallel p(z))$$

- KL term enforces *local* prior matching for every input
- Often leads to posterior collapse

- **MMD-VAE Loss**

$$\mathcal{L}_{\text{MMD-VAE}} = \mathbb{E}_{q(z|x)}[\log p(x|z)] - \lambda \text{MMD}(q(z), p(z))$$

- MMD enforces *global* matching of the aggregated posterior
- Encourages higher mutual information and better latent usage

Why MMD-VAE? Advantages over Vanilla VAE

- **Avoids Posterior Collapse**

- Vanilla VAE:

$$q(z | x) \rightarrow \mathcal{N}(0, I), \quad \forall x$$

- MMD-VAE:

$$\mathbb{E}_x [q(z | x)] \approx \mathcal{N}(0, I)$$

- Prior matching is enforced *globally* (in expectation), not per-sample
- Encourages higher mutual information between x and z

- **Better Balance Between Reconstruction and Regularization**

- In Vanilla VAE, KL term often becomes negligible compared to reconstruction loss
- MMD penalty remains influential throughout training
- Leads to more meaningful latent representations

Mutual Information: ConvVAE vs MMD-VAE

- Experiment conducted using a **simple (non-complex) encoder-decoder architecture**
- Mutual Information (MI) measures dependency between input x and latent code z

Model	Mutual Information ↑
ConvVAE	5.41
MMD-VAE	5.54

MMD-VAE preserves slightly higher mutual information, indicating stronger latent utilization even with a simple architecture.

MMD-VAE: Overall Performance on CIFAR-10

- Evaluated using sample quality, diversity, and latent utilization metrics


Metric	Value
FID ↓	116
IS ↑	4.105
Mutual Information ↑	5.54

MMD-VAE achieves strong generative performance while maintaining high latent information utilization.

Generated Samples (MMD Vae)




Reconstruction Samples (MMD-VAE)



VQ-VAE

- Discrete latent space with vector quantization
- Learns a codebook of embeddings
- Effective for high-quality reconstructions

VQ-VAE Architecture



VQ-VAE Loss Function

- VQ-VAE uses a **discrete latent space** with vector quantization

$$\mathcal{L}_{\text{VQ-VAE}} = \underbrace{\|x - \hat{x}\|^2}_{\text{Reconstruction}} + \underbrace{\|\text{sg}[z_e(x)] - e_k\|^2}_{\text{Codebook Loss}} + \underbrace{\beta \|z_e(x) - \text{sg}[e_k]\|^2}_{\text{Commitment Loss}}$$

- $z_e(x)$: encoder output
- e_k : nearest codebook embedding
- $\text{sg}[\cdot]$: stop-gradient operator

Why VQ-VAE? Key Advantages

- **Sharper Reconstructions**

- Vanilla VAE samples multiple latents:

$$z_1, z_2 \sim q(z | x)$$

- Decoder must map different z values to the same x
- Averaging effect leads to blurry reconstructions
- VQ-VAE maps x to a single discrete code e_k
- No averaging \Rightarrow sharper images

- **No KL Term & No Posterior Collapse**

- No prior matching via KL divergence
- Discrete bottleneck enforces latent usage
- Encoder cannot ignore the latent representation

VQ-VAE: Generative Process

- Encoder maps input image to discrete latent codes:

$$x \xrightarrow{\text{Encoder}} z_e(x) \xrightarrow{\text{VQ}} e_k$$

- A **secondary autoregressive model** (e.g., PixelCNN) is trained on the discrete latents:

$$p(e) = \prod_i p(e_i | e_{<i})$$

- Sampling pipeline:

$$e \sim \text{PixelCNN} \rightarrow \hat{x} = \text{Decoder}(e)$$


- Disadvantage: PixelCNN Bottleneck**

- Requires training an additional complex model
- Autoregressive sampling is **slow and sequential**
- Sampling speed limited by PixelCNN, not the decoder


VQ-VAE: Quantitative Results on CIFAR-10

Metric	Value
FID ↓	84.45
IS ↑	4.448
Mutual Information ↑	5.36
Reconstruction Loss ↓	10.72

Generated Samples (VQ-VAE)



Reconstruction Samples (VQ-VAE)



Summary: VAE-Based Models on CIFAR-10

Model	FID ↓	IS ↑	MI ↑	Recon. Loss ↓
ConvVAE ($\beta = 0.1$)	129	3.66	5.53	31.66
MMD-VAE	116	4.105	5.54	7.54
VQ-VAE	84.45	4.448	5.36	10.72

Progressive improvement in sample quality is observed from ConvVAE to VQ-VAE.


- Noise Conditional Score Network
- Learns gradients of the data distribution
- Sampling via Langevin dynamics
- Unet-type network

NCSN: Quantitative Results on CIFAR-10

Metric	Value
FID ↓	86.43
IS ↑	4.12

NCSN achieves competitive sample quality without an explicit latent representation.

NCSN: FID Over Training Epochs



DDPM Architecture and Training Setup

- **Backbone: U-Net**

- Residual blocks with Group Normalization
- Self-attention at 16×16 resolution
- Sinusoidal position embeddings for timestep encoding
- Approximately **37M parameters** (CIFAR-10 configuration)

- **Diffusion Process (Following the DDPM Paper)**

- Number of timesteps: $T = 1000$
- Linear noise schedule:

$$\beta_1 = 10^{-4} \rightarrow \beta_T = 0.02$$

- Dropout rate: 0.1
- Exponential Moving Average (EMA) of parameters with decay 0.9999

DDPM Training Algorithm

- ➊ Sample a clean image: $x_0 \sim p_{\text{data}}(x)$
- ➋ Sample timestep: $t \sim \mathcal{U}\{1, \dots, T\}$
- ➌ Sample noise: $\varepsilon \sim \mathcal{N}(0, I)$
- ➍ Corrupt the image:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon$$

- ➎ Predict noise: $\varepsilon_\theta(x_t, t)$

- ➏ Compute loss:

$$\mathcal{L} = \|\varepsilon - \varepsilon_\theta(x_t, t)\|^2$$

- ➐ Update model parameters θ via gradient descent

DDPM Sampling Algorithm (Reverse Diffusion)

- ① Initialize with pure noise:

$$x_T \sim \mathcal{N}(0, I)$$

- ② For $t = T, \dots, 1$:

- Predict noise:

$$\varepsilon_\theta(x_t, t)$$

- Compute the mean:


$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right)$$

- Sample:

$$x_{t-1} = \mu_\theta + \sqrt{\tilde{\beta}_t} z, \quad z \sim \begin{cases} \mathcal{N}(0, I), & t > 1 \\ 0, & t = 1 \end{cases}$$

- ③ Return generated sample: x_0

DDPM Denoising Process



DDPM: Model Components

- **Noise Schedule**
 - `get_beta_schedule()`: constructs β_t (linear or cosine)
- **Timestep Encoding**
 - SinusoidalPositionEmbeddings
 - Encodes diffusion timestep t (Transformer-style)
- **Residual Block**
 - Group Normalization
 - Time embedding injection
 - Dropout
- **AttentionBlock**
 - Multi-head self-attention
- **U-Net Architecture**
 - Downsampling path with attention layers
 - Middle blocks at lowest resolution
 - Upsampling path with skip connections

DDPM: Main API

- **DDPM**

- Main diffusion model wrapper

- `q_sample()`

- Forward diffusion process
 - Adds noise to clean images:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon$$

- `p_losses()`

- Computes training objective
 - MSE between true and predicted noise

- `p_sample()`

- One-step reverse diffusion (denoising)
 - Computes x_{t-1} from x_t

- `sample()`

- Full generative procedure
 - Iteratively denoises from $x_T \rightarrow x_0$

Generated Samples (DDPM)



Conclusion

- Compared four generative paradigms on CIFAR-10
- Trade-off between sample quality, diversity, and training cost
- Future work: hybrid models, larger datasets

Thank You

Questions?